

# DataLunch Statistical Power

Ian Dworkin

23 Jul 2025

# Points to think about before starting your power analysis




Figure 1. Confidence intervals for five different environmental scenarios.

Figure 1: Effects\_Not\_Pvalues

Figure 1 from Fox (2001)

# Points to think about before starting your power analysis


- ▶ Power analysis to maximize precision of your quantities of interest
- ▶ Lakens (2022) is a good starting place for considering how to determine appropriate/sufficient sample sizes (including a Shiny app!)

the traditional “4 possible outcomes of a statistical test”


	Reject Null	Accept Null
Null True	Type I error, $\alpha$	Correct, $1 - \alpha$
Null False	Correct, $1 - \beta$	Type II error, $\beta$

- ▶  $(1 - \beta)$  is **power**, probability of detecting a true difference.
- ▶  $(1 - \alpha)$  is **confidence**, probability of correctly accepting null.

# Critical value for a $t$ distribution, for a one tailed test



# Critical value for a $t$ distribution, for a two tailed test



## Keep in mind

These kinds of dichotomies lead you to an *“Is there an effect?”* thinking.

Instead you should ask *“What is the effect?”* and for a power analysis, *“What precision of the effect do I want, given the resources I have?”*

## What does a p-value tell you?

- ▶ Say you conduct an analysis on two different data sets, in the first  $p = 0.05$ , the second test has  $p = 0.001$ .



## What does a p-value tell you?

- ▶ Say you conduct an analysis on two different data sets, in the first  $p = 0.05$ , the second test has  $p = 0.001$ .
- ▶ Does this mean the second test has a bigger effect? That the statistical model accounts for more variation (i.e. higher  $R^2$ )?

## What does a p-value tell you?

- ▶ Say you conduct an analysis on two different data sets, in the first  $p = 0.05$ , the second test has  $p = 0.001$ .
- ▶ Does this mean the second test has a bigger effect? That the statistical model accounts for more variation (i.e. higher  $R^2$ )?
- ▶ Not necessarily. The magnitude of an effect could be similar, and the sample sizes differ (the second data set being much larger).

## What does a p-value tell you?

- ▶ Say you conduct an analysis on two different data sets, in the first  $p = 0.05$ , the second test has  $p = 0.001$ .
- ▶ Does this mean the second test has a bigger effect? That the statistical model accounts for more variation (i.e. higher  $R^2$ )?
- ▶ Not necessarily. The magnitude of an effect could be similar, and the sample sizes differ (the second data set being much larger).
- ▶ It could also be that there is less variability in the second data set.

## What does a p-value tell you?


- ▶ Say you conduct an analysis on two different data sets, in the first  $p = 0.05$ , the second test has  $p = 0.001$ .
- ▶ Does this mean the second test has a bigger effect? That the statistical model accounts for more variation (i.e. higher  $R^2$ )?
- ▶ Not necessarily. The magnitude of an effect could be similar, and the sample sizes differ (the second data set being much larger).
- ▶ It could also be that there is less variability in the second data set.
- ▶ However it could also be that there is a difference in these. You need to examine (and report) all three whenever possible (include confidence intervals on estimates).

## Let's compare these three data sets




- ▶ We are examining the relationship between body size and fitness.

# Let's compare these three data sets

- Is there a relationship?



## Is there a relationship?



- In fact they all have the same relationship  $fitness \sim N(2 + 0.5 * size, \sigma = 2)$ , and only differ in sample size.

# Statistical Power Analysis in R

- ▶ Most statistical software packages provide functions for simple power analyses



# Statistical Power Analysis in R

- ▶ Most statistical software packages provide functions for simple power analyses
- ▶ R has lots of useful packages

## Getting a critical value (for $t$ distribution)

- ▶ Let's start by seeing how we get critical values for a  $t$  distribution

## Getting a critical value (for $t$ distribution)

- ▶ Let's start by seeing how we get critical values for a  $t$  distribution
- ▶ Assume we have a sample size of  $n = 25$ ,  $\alpha = 0.05$  for a two-tailed distribution.

## Getting a critical value (for $t$ distribution)

- ▶ Let's start by seeing how we get critical values for a  $t$  distribution
- ▶ Assume we have a sample size of  $n = 25$ ,  $\alpha = 0.05$  for a two-tailed distribution.
- ▶ We can use the `qt()` function for the  $t$  distribution

## Getting a critical value (for $t$ distribution)

- ▶ Let's start by seeing how we get critical values for a  $t$  distribution
- ▶ Assume we have a sample size of  $n = 25$ ,  $\alpha = 0.05$  for a two-tailed distribution.
- ▶ We can use the `qt()` function for the  $t$  distribution

```
qt(p = 0.975, df = 24)
```

```
## [1] 2.063899
```

- ▶ Why do we have  $df = 24$ , not 25?

## Getting a critical value (for $t$ distribution)

- ▶ Let's start by seeing how we get critical values for a  $t$  distribution
- ▶ Assume we have a sample size of  $n = 25$ ,  $\alpha = 0.05$  for a two-tailed distribution.
- ▶ We can use the `qt()` function for the  $t$  distribution

```
qt(p = 0.975, df = 24)
```

```
## [1] 2.063899
```

- ▶ Why do we have  $df = 24$ , not 25?
- ▶ Why is  $p = 0.975$ , not 0.95 (with  $\alpha = 0.05$ )?

## How does the critical value change with sample size?

- ▶ We can make a plot looking at this across a range of sample sizes.

## How does the critical value change with sample size?


- We can make a plot looking at this across a range of sample sizes.

```
curve(qt(p = 0.975,df = x), 2, 25,  
      col = "red", lwd = 3, cex.lab = 2,  
      main = "Critical values of t for different sample sizes",  
      xlab = "Sample Size",  
      ylab = expression(paste("t value (two tailed)", alpha)))
```



## How does the critical value change with sample size?

- We can make a plot looking at this across a range of sample sizes.



## Critical values for other distributions

- ▶ There are other distributions we can use:

## Critical values for other distributions

- ▶ There are other distributions we can use:
- ▶ `qf()` for the  $F$  distribution, `qchisq()` for  $\chi^2$  etc..

## tools in R

- ▶ Many libraries in R to use, depending on purposes.

## tools in R

- ▶ Many libraries in R to use, depending on purposes.
- ▶ in base R, there is `power.t.test()`, `power.anova.test()`, `power.prop.test()`

## tools in R

- ▶ Many libraries in R to use, depending on purposes.
- ▶ in base R, there is `power.t.test()`, `power.anova.test()`, `power.prop.test()`
- ▶ `pwr` is an R package that does many simple types of statistical models (proportions, 1-way ANOVA, regression,  $\chi^2$ , glm)

## tools in R

- ▶ Many libraries in R to use, depending on purposes.
- ▶ in base R, there is `power.t.test()`, `power.anova.test()`, `power.prop.test()`
- ▶ `pwr` is an R package that does many simple types of statistical models (proportions, 1-way ANOVA, regression,  $\chi^2$ , glm)
- ▶ See the CRAN TASK VIEWS on experimental design for a list (and description) of more.

## tools in R

- ▶ Many libraries in R to use, depending on purposes.
- ▶ in base R, there is `power.t.test()`, `power.anova.test()`, `power.prop.test()`
- ▶ `pwr` is an R package that does many simple types of statistical models (proportions, 1-way ANOVA, regression,  $\chi^2$ , glm)
- ▶ See the CRAN TASK VIEWS on experimental design for a list (and description) of more.
- ▶ See this draft task view for power



## tools in R

- ▶ Many libraries in R to use, depending on purposes.
- ▶ in base R, there is `power.t.test()`, `power.anova.test()`, `power.prop.test()`
- ▶ `pwr` is an R package that does many simple types of statistical models (proportions, 1-way ANOVA, regression,  $\chi^2$ , glm)
- ▶ See the CRAN TASK VIEWS on experimental design for a list (and description) of more.
- ▶ See this draft task view for power
- ▶ I will show just a couple here.

## Some of the functions in base R

```
## [1] "power"
```

```
"power.anova.test" "power.prop.test"
```

power.t.test

- ▶ What goes into a  $t$ -test?

## power.t.test

- ▶ What goes into a  $t$ -test?



$$\frac{\bar{X}_A - \bar{X}_B}{\hat{\sigma} \frac{1}{\sqrt{n}}}$$

## power.t.test

- ▶ What goes into a  $t$ -test?



$$\frac{\bar{x}_A - \bar{x}_B}{\hat{\sigma} \frac{1}{\sqrt{n}}}$$

- ▶  $\bar{x}_A$  is the mean for group  $A$ ,  $\bar{x}_B$  for  $B$

## power.t.test

- ▶ What goes into a  $t$ -test?



$$\frac{\bar{x}_A - \bar{x}_B}{\hat{\sigma} \frac{1}{\sqrt{n}}}$$

- ▶  $\bar{x}_A$  is the mean for group  $A$ ,  $\bar{x}_B$  for  $B$
- ▶ The denominator is just the *pooled standard error of the mean*

## power.t.test

- ▶ What goes into a  $t$ -test?



$$\frac{\bar{x}_A - \bar{x}_B}{\hat{\sigma} \frac{1}{\sqrt{n}}}$$

- ▶  $\bar{x}_A$  is the mean for group  $A$ ,  $\bar{x}_B$  for  $B$
- ▶ The denominator is just the *pooled standard error of the mean*
- ▶ So we see that there are 4 critical things:

## power.t.test

- ▶ What goes into a  $t$ -test?



$$\frac{\bar{x}_A - \bar{x}_B}{\hat{\sigma} \frac{1}{\sqrt{n}}}$$

- ▶  $\bar{x}_A$  is the mean for group  $A$ ,  $\bar{x}_B$  for  $B$
- ▶ The denominator is just the *pooled standard error of the mean*
- ▶ So we see that there are 4 critical things:
- ▶  $\alpha$ , the difference between means  $\Delta = \bar{x}_A - \bar{x}_B$ ,  $n$  and  $\hat{\sigma}$



## power.t.test

```
pwr_t_check <- power.t.test(delta = 0.5, sd = 2,  
                             sig.level = 0.05,  
                             power = 0.8)  
  
pwr_t_check
```

```
##  
##      Two-sample t test power calculation  
##  
##              n = 252.1281  
##            delta = 0.5  
##              sd = 2  
##    sig.level = 0.05  
##            power = 0.8  
##    alternative = two.sided  
##  
## NOTE: n is number in *each* group
```

```
str(pwr_t_check)
```

## power.t.test

- ▶ what sample sizes we would need for a range of differences,  $\Delta$ , on the interval  $[0.1, 0.5]$ .

## power.t.test

- ▶ what sample sizes we would need for a range of differences,  $\Delta$ , on the interval  $[0.1, 0.5]$ .
- ▶  $(1 - \beta) = 0.8$ ,  $\hat{\sigma} = 2$ ,  $\alpha = 0.05$

## power.t.test

- ▶  $\Delta = 0.5$ ,  $\hat{\sigma} = 2$ ,  $\alpha = 0.05$

```
delta_vals = seq(from = 0.1, to = 0.5, by = 0.01)
delta_vals
```

```
## [1] 0.10 0.11 0.12 0.13 0.14 0.15 0.16 0.17 0.18 0.19 0
## [16] 0.25 0.26 0.27 0.28 0.29 0.30 0.31 0.32 0.33 0.34 0
## [31] 0.40 0.41 0.42 0.43 0.44 0.45 0.46 0.47 0.48 0.49 0
```

- ▶ This creates a vector from 0.1 - 0.5

power.t.test

```
pow.test <- function(x) {  
  pow2 <- power.t.test(delta = x, sd = 2,  
                        sig.level = 0.05,  
                        power = 0.8) # We only allow delta  
  return(pow2$n) # Pull out the sample size we need  
}
```

power.t.test

```
power.n <- sapply(delta_vals, pow.test)
```


- ▶ This uses one of the \*apply functions to repeat the function `pow.test` for each element of the vector `delta_vals`.

power.t.test

```
power.n <- sapply(delta_vals, pow.test)
```

- ▶ This uses one of the \*apply functions to repeat the function `pow.test` for each element of the vector `delta_vals`.
- ▶ Thus for each value in the vector `delta_vals` (from 0.1 to 0.5), it puts this value into `pow.test()` and then returns the estimated `n` (# of observations needed to achieve this power).

# power.t.test





Similarly, there are functions in base R for 1-way ANOVA

`power.anova.test` example

## More complex power analyses

- ▶ `pwr` has many useful functions for experimental designs of simple to moderate complexity.
- ▶ `pwrss` does as well, and can generate some very helpful figures to help understand
- ▶ If you are designing experiments and you think it is likely you are going to use mixed models, the `simr` is a good choice to learn (relatively straightforward)
- ▶ `EMSS` has useful sample size calculators.

## roll your own with Monte Carlo simulations

- ▶ It is relatively straightforward to put this in a for loop to generate power analyses for more complicated designs/analyses

## roll your own with Monte Carlo simulations

- ▶ It is relatively straightforward to put this in a for loop to generate power analyses for more complicated designs/analyses
- ▶ Learning how to do simple *Monte Carlo* simulations gives you lots of flexibility


## roll your own with Monte Carlo simulations

- ▶ It is relatively straightforward to put this in a `for` loop to generate power analyses for more complicated designs/analyses
- ▶ Learning how to do simple *Monte Carlo* simulations gives you lots of flexibility
- ▶ I have posted a series of screencasts on YouTube, starting here that will teach you the basics.


# Monte carlo power analysis example

- ▶ R code is hidden (but you can see it with the .Rmd file)

# Plotting results from a power analysis




## Plotting results from a power analysis





# Plotting results from a power analysis



# References

- Fox, David R. 2001. "Environmental Power Analysis – a New Perspective." *Environmetrics* 12 (5): 437–49.  
<https://doi.org/10.1002/env.470>.
- Lakens, Daniël. 2022. "Sample Size Justification." *Collabra: Psychology* 8 (1): 33267.  
<https://doi.org/10.1525/collabra.33267>.