

DataLunch Statistical Power

Ian Dworkin

14 Jul 2023

points to think about before starting your power analysis




Figure 1. Confidence intervals for five different environmental scenarios.

Figure 1: Effects_Not_Pvalues

From Fox 2001 (DOI:10.1002/env.470)

points to think about before starting your power analysis


- ▶ Power analysis to maximize precision of your quantities of interest
- ▶ This paper by Daniël Lakens is a good starting place for considering how to determine sample sizes that will be sufficient for your needs.

the traditional “4 possible outcomes of a statistical test”


	Reject Null	Accept Null
Null True	Type I error, α	Correct, $1 - \alpha$
Null False	Correct, $1 - \beta$	Type II error, β

- ▶ $(1 - \beta)$ is power, probability of detecting a true difference.
- ▶ $(1 - \alpha)$ is confidence, probability of correctly accepting null.

Critical value for a t distribution, for a one tailed test



Critical value for a t distribution, for a two tailed test



Keep in mind

These kinds of dichotomies lead you to an *“Is there an effect?”* thinking.

Instead you should ask *“What is the effect?”* and for a power analysis, *“What precision of the effect do I want, given the resources I have?”*

What does a p-value tell you?

- ▶ Say you conduct an analysis on two different data sets, in the first $p = 0.05$, the second test has $p = 0.001$.

What does a p-value tell you?

- ▶ Say you conduct an analysis on two different data sets, in the first $p = 0.05$, the second test has $p = 0.001$.
- ▶ Does this mean the second test has a bigger effect? That the statistical model accounts for more variation (i.e. higher R^2)?

What does a p-value tell you?

- ▶ Say you conduct an analysis on two different data sets, in the first $p = 0.05$, the second test has $p = 0.001$.
- ▶ Does this mean the second test has a bigger effect? That the statistical model accounts for more variation (i.e. higher R^2)?
- ▶ Not necessarily. The magnitude of an effect could be similar, and the sample sizes differ (the second data set being much larger).

What does a p-value tell you?

- ▶ Say you conduct an analysis on two different data sets, in the first $p = 0.05$, the second test has $p = 0.001$.
- ▶ Does this mean the second test has a bigger effect? That the statistical model accounts for more variation (i.e. higher R^2)?
- ▶ Not necessarily. The magnitude of an effect could be similar, and the sample sizes differ (the second data set being much larger).
- ▶ It could also be that there is less variability in the second data set.

What does a p-value tell you?


- ▶ Say you conduct an analysis on two different data sets, in the first $p = 0.05$, the second test has $p = 0.001$.
- ▶ Does this mean the second test has a bigger effect? That the statistical model accounts for more variation (i.e. higher R^2)?
- ▶ Not necessarily. The magnitude of an effect could be similar, and the sample sizes differ (the second data set being much larger).
- ▶ It could also be that there is less variability in the second data set.
- ▶ However it could also be that there is a difference in these. You need to examine (and report) all three whenever possible (include confidence intervals on estimates).

Let's compare these three data sets




- ▶ We are examining the relationship between body size and fitness.

Let's compare these three data sets

- Is there a relationship?



Is there a relationship?



- In fact they all have the same relationship $fitness \sim N(2 + 0.5 * size, \sigma = 2)$, and only differ in sample

Statistical Power Analysis in R

- ▶ Most statistical software packages provide functions for simple power analyses

Statistical Power Analysis in R

- ▶ Most statistical software packages provide functions for simple power analyses
- ▶ In R there are many libraries one can use

Getting a critical value (for t distribution)

- ▶ Let's start by seeing how we get critical values for a t distribution

Getting a critical value (for t distribution)

- ▶ Let's start by seeing how we get critical values for a t distribution
- ▶ Assume we have a sample size of $n = 25$, $\alpha = 0.05$ for a two-tailed distribution.

Getting a critical value (for t distribution)

- ▶ Let's start by seeing how we get critical values for a t distribution
- ▶ Assume we have a sample size of $n = 25$, $\alpha = 0.05$ for a two-tailed distribution.
- ▶ We can use the `qt()` function for the t distribution

Getting a critical value (for t distribution)

- ▶ Let's start by seeing how we get critical values for a t distribution
- ▶ Assume we have a sample size of $n = 25$, $\alpha = 0.05$ for a two-tailed distribution.
- ▶ We can use the `qt()` function for the t distribution

```
qt(p = 0.975, df = 24)
```

```
## [1] 2.06
```

- ▶ Why do we have $df = 24$, not 25?

Getting a critical value (for t distribution)

- ▶ Let's start by seeing how we get critical values for a t distribution
- ▶ Assume we have a sample size of $n = 25$, $\alpha = 0.05$ for a two-tailed distribution.
- ▶ We can use the `qt()` function for the t distribution

```
qt(p = 0.975, df = 24)
```

```
## [1] 2.06
```

- ▶ Why do we have $df = 24$, not 25?
- ▶ Why is $p = 0.975$, not 0.95 (with $\alpha = 0.05$)?

How does the critical value change with sample size?

- ▶ We can make a plot looking at this across a range of sample sizes.

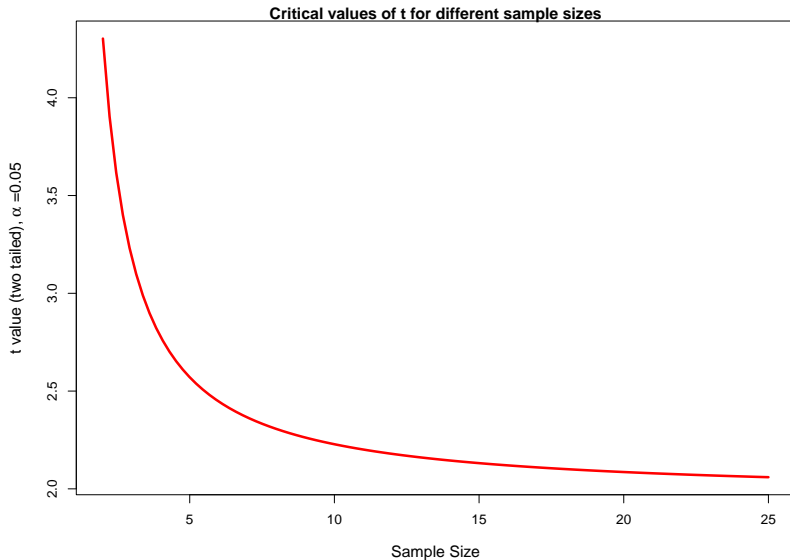
How does the critical value change with sample size?

- ▶ We can make a plot looking at this across a range of sample sizes.

```
curve(qt(p = 0.975, df = x), 2, 25,  
      col = "red", lwd = 3, cex.lab = 2,  
      main = "Critical values of t for different sample sizes",  
      xlab = "Sample Size",  
      ylab = expression(paste("t value (two tailed)", alpha)))
```


How does the critical value change with sample size?

- ▶ We can make a plot looking at this across a range of sample sizes.



Critical values for other distributions

- ▶ There are other distributions we can use:

Critical values for other distributions

- ▶ There are other distributions we can use:
- ▶ `qf()` for the F distribution, `qchisq()` for χ^2 etc..

tools in R

- ▶ Many libraries in R to use, depending on purposes.

tools in R

- ▶ Many libraries in R to use, depending on purposes.
- ▶ in base R, there is `power.t.test()`, `power.anova.test()`, `power.prop.test()`

tools in R

- ▶ Many libraries in R to use, depending on purposes.
- ▶ in base R, there is `power.t.test()`, `power.anova.test()`, `power.prop.test()`
- ▶ `pwr` is an R package that does many simple types of statistical models (proportions, 1-way ANOVA, regression, χ^2 , glm)

tools in R

- ▶ Many libraries in R to use, depending on purposes.
- ▶ in base R, there is `power.t.test()`, `power.anova.test()`, `power.prop.test()`
- ▶ `pwr` is an R package that does many simple types of statistical models (proportions, 1-way ANOVA, regression, χ^2 , glm)
- ▶ See the CRAN TASK VIEWS on experimental design for a list (and description) of more.

tools in R

- ▶ Many libraries in R to use, depending on purposes.
- ▶ in base R, there is `power.t.test()`, `power.anova.test()`, `power.prop.test()`
- ▶ `pwr` is an R package that does many simple types of statistical models (proportions, 1-way ANOVA, regression, χ^2 , glm)
- ▶ See the CRAN TASK VIEWS on experimental design for a list (and description) of more.
- ▶ See this draft task view for power

tools in R

- ▶ Many libraries in R to use, depending on purposes.
- ▶ in base R, there is `power.t.test()`, `power.anova.test()`, `power.prop.test()`
- ▶ `pwr` is an R package that does many simple types of statistical models (proportions, 1-way ANOVA, regression, χ^2 , glm)
- ▶ See the CRAN TASK VIEWS on experimental design for a list (and description) of more.
- ▶ See this draft task view for power
- ▶ I will show just a couple here.

Some of the functions in base R

```
## [1] "power"
```

```
"power.anova.test" "power.prop.test"
```

power.t.test

- ▶ What goes into a t -test?

power.t.test

- ▶ What goes into a t -test?



$$\frac{\bar{x}_A - \bar{x}_B}{\hat{\sigma} \frac{1}{\sqrt{n}}}$$

power.t.test

- ▶ What goes into a t -test?



$$\frac{\bar{x}_A - \bar{x}_B}{\hat{\sigma} \frac{1}{\sqrt{n}}}$$

- ▶ \bar{x}_A is the mean for group A , \bar{x}_B for B

power.t.test

- ▶ What goes into a t -test?



$$\frac{\bar{x}_A - \bar{x}_B}{\hat{\sigma} \frac{1}{\sqrt{n}}}$$

- ▶ \bar{x}_A is the mean for group A , \bar{x}_B for B
- ▶ The denominator is just the *pooled standard error of the mean*

power.t.test

- ▶ What goes into a t -test?



$$\frac{\bar{x}_A - \bar{x}_B}{\hat{\sigma} \frac{1}{\sqrt{n}}}$$

- ▶ \bar{x}_A is the mean for group A , \bar{x}_B for B
- ▶ The denominator is just the *pooled standard error of the mean*
- ▶ So we see that there are 4 critical things:

power.t.test

- ▶ What goes into a t -test?



$$\frac{\bar{x}_A - \bar{x}_B}{\hat{\sigma} \frac{1}{\sqrt{n}}}$$

- ▶ \bar{x}_A is the mean for group A , \bar{x}_B for B
- ▶ The denominator is just the *pooled standard error of the mean*
- ▶ So we see that there are 4 critical things:
- ▶ α , the difference between means $\Delta = \bar{x}_A - \bar{x}_B$, n and $\hat{\sigma}$

power.t.test

```
pwr_t_check <- power.t.test(delta = 0.5, sd = 2,  
                             sig.level = 0.05, power = 0.8)
```

```
pwr_t_check
```

```
##  
##      Two-sample t test power calculation  
##  
##              n = 252  
##          delta = 0.5  
##            sd = 2  
##    sig.level = 0.05  
##          power = 0.8  
## alternative = two.sided  
##  
## NOTE: n is number in *each* group
```

```
str(pwr_t_check)
```

power.t.test

- ▶ what sample sizes we would need for a range of differences, $\Delta = (0.1 - -0.5)$.

power.t.test

- ▶ what sample sizes we would need for a range of differences, $\Delta = (0.1 - -0.5)$.
- ▶ $(1 - \beta) = 0.8, \hat{\sigma} = 2, \alpha = 0.05$

power.t.test

► $\Delta = 0.5$, $\hat{\sigma} = 2$, $\alpha = 0.05$

```
delta_vals = seq(from = 0.1, to = 0.5, by = 0.01)
delta_vals
```

```
## [1] 0.10 0.11 0.12 0.13 0.14 0.15 0.16 0.17 0.18 0.19 0
## [16] 0.25 0.26 0.27 0.28 0.29 0.30 0.31 0.32 0.33 0.34 0
## [31] 0.40 0.41 0.42 0.43 0.44 0.45 0.46 0.47 0.48 0.49 0
```

► This creates a vector from 0.1 - 0.5

power.t.test

```
pow.test <- function(x){  
  pow2 <- power.t.test(delta = x, sd = 2,  
                        sig.level = 0.05, power = 0.8) # We  
  return(pow2$n) # This pulls out the sample size we need  
}
```

power.t.test

```
power.n <- sapply(delta_vals, pow.test)
```


- ▶ This just uses one of the apply functions to repeat the function `pow.test` for each element of the vector “delta_vals”.

power.t.test

```
power.n <- sapply(delta_vals, pow.test)
```

- ▶ This just uses one of the apply functions to repeat the function `pow.test` for each element of the vector “delta_vals”.
- ▶ Thus for each value in the vector “delta_vals” (from 0.1 to 0.5), it inputs this value into `pow.test()` and then returns the estimated `n` (# of observations needed to achieve this power).

power.t.test



Similarly, there are functions in base R for 1-way ANOVA

`power.anova.test` example

More complex power analyses

- ▶ `pwr` has many useful functions for experimental designs of simple to moderate complexity.
- ▶ `pwrss` does as well, and can generate some very helpful figures to help understand
- ▶ If you are designing experiments and you think it is likely you are going to use mixed models, the `simr` is a good choice to learn (relatively straightforward)
- ▶ `EMSS` has useful sample size calculators.

role your own with monte carlo simulations

- ▶ It is relatively straightforward to loop this and generate more complex power analyses.

role your own with monte carlo simulations

- ▶ It is relatively straightforward to loop this and generate more complex power analyses.
- ▶ Learning how to do simple *Monte Carlo* simulations can give you a lot of flexibility to do this.

role your own with monte carlo simulations

- ▶ It is relatively straightforward to loop this and generate more complex power analyses.
- ▶ Learning how to do simple *Monte Carlo* simulations can give you a lot of flexibility to do this.
- ▶ I have posted a series of screencasts on youtube, starting here that will teach you the basics.

Monte carlo power analysis example

- ▶ R code is hidden (but you can see it with the .Rmd file)

Plotting results from a power analysis

Plotting results from a power analysis

Plotting results from a power analysis